

HOWTO-template for Big HOWTOs

Stein Gjoen

sgjoen@nyx.net

Gregory Leblanc

gleblanc@cu-portland.edu

Greg Ferguson

gferg@sgi.com

Revision History

Revision 0.06 2001-01-08 Revised by: gjf
Changed date format (YYYY-MM-DD)

This is a fully working template for big HOWTOs. The source contains fully described slots to make a convenient framework for you to fill in for making your own HOWTO; suggesting some names, conventions and contents for the chapters.

SGML source (big-howto-template.sgml) for this document is available and should be consulted while reading.

1. Introduction

For various reasons this brand new release is codenamed the *release* release.

New code names will appear as per industry standard guidelines to emphasize the state-of-the-art-ness of

this document.

This document was written when I read a feedback asking for a template to fill in to make new HOWTOs. This template was initially made by extracting the skeletal structure of the Multi Disk HOWTO which is a rather large HOWTO. It then went through extensive editing.

Stating the background is a simple way to getting started writing the intro.

First of all we need a bit of legalese. Recent development shows it is quite important.

1.1. Copyright Information

This document is copyrighted (c) 2000 Stein Gjoen and is distributed under the terms of the Linux Documentation Project (LDP) license, stated below. *Replace with your name, or supply a new license, when you use this skeleton for a new HOWTO.*

Unless otherwise stated, Linux HOWTO documents are copyrighted by their respective authors. Linux HOWTO documents may be reproduced and distributed in whole or in part, in any medium physical or electronic, as long as this copyright notice is retained on all copies. Commercial redistribution is allowed and encouraged; however, the author would like to be notified of any such distributions.

All translations, derivative works, or aggregate works incorporating any Linux HOWTO documents must be covered under this copyright notice. That is, you may not produce a derivative work from a HOWTO and impose additional restrictions on its distribution. Exceptions to these rules may be granted under certain conditions; please contact the Linux HOWTO coordinator at the address given below.

In short, we wish to promote dissemination of this information through as many channels as possible. However, we do wish to retain copyright on the HOWTO documents, and would like to be notified of any plans to redistribute the HOWTOs.

If you have any questions, please contact <linux-howto@metalab.unc.edu>

1.2. Disclaimer

No liability for the contents of this documents can be accepted. Use the concepts, examples and other content at your own risk. As this is a new edition of this document, there may be errors and inaccuracies, that may of course be damaging to your system. Proceed with caution, and although this is highly unlikely, the author(s) do not take any responsibility for that.

All copyrights are held by their by their respective owners, unless specifically noted otherwise. Use of a term in this document should not be regarded as affecting the validity of any trademark or service mark.

Naming of particular products or brands should not be seen as endorsements.

You are strongly recommended to take a backup of your system before major installation and backups at regular intervals.

1.3. New Versions

This is where you make a summary of what is new. When a HOWTO exceeds 20 pages it takes more than a casual read to find the updates. This is where you help your readers with that, alerting them to specific and important updates to the document.

This is the initial release.

Tell people where the document home page is so the very newest release could be found in case of problems with the main Linux Documentation Project (<http://www.linuxdoc.org/>) homepage.

The following is a sample from the Multi Disk HOWTO:

The latest version number of this document can be gleaned from my plan entry if you `finger` (<http://www.cs.indiana.edu/finger/nox.nyx.net/sgjoen>) my Nyx account.

If you have the capability, it would be nice to make the HOWTO available in a number of formats.

The newest version of this HOWTO will always be made available on my website, in a variety of formats:

- HTML (<http://www.nyx.net/~sgjoen/disk.html>).
- plain text (<http://www.nyx.net/~sgjoen/disk.txt>).
- compressed postscript (US letter format) (<http://www.nyx.net/~sgjoen/disk-US.ps.gz>).
- SGML source (<http://www.nyx.net/~sgjoen/disk.sgml>).

Note that paper sizes vary in the world, A4 and US letter differ significantly. You might also wish to consider using the *universal format* (8.27x11in; 210x279mm).

1.4. Credits

It is always nice to acknowledge people who help you with input; it is also regarded by many as important in the Linux world new economy.

In this version I have the pleasure of acknowledging:

<name (at) site.org>

Please scramble the addresses so email harvesters cannot get addresses from your HOWTO and then spam people. That has happened in the past.

Somecompany is acknowledged for sending me documentation on their gizmos as well as permission to quote from the material. These quotes have been approved before appearing here and will be clearly labeled.

1.5. Feedback

Feedback is most certainly welcome for this document. Without your submissions and input, this document wouldn't exist. Please send your additions, comments and criticisms to the following email address : <sgjoen@nyx.net>.

1.6. Translations

Not everyone speaks English, pointers to translations are nice. Also your translators tend to give very important inputs.

- German Translation (<http://linuxdoc.org/>) by <someone (at) somewhere.de>
- French Translation (<http://linuxdoc.org/>) by <someone (at) somewhere.fr>
- Italian Translation (<http://linuxdoc.org/>) by <someone (at) somewhere.it>

2. Samples

This section gives some simple SGML examples you could use. Read the SGML source to see how it was done.

Further information and examples can be obtained from the publication DocBook: The Definitive Guide (<http://docbook.org/tdg/html/>). Written by *Norman Walsh* and *Leonard Mueller*; 1st Edition, October 1999.

2.1. Lists

Lists are used frequently, and are available in a number of formats shown below.

A list in which each entry is marked with a bullet or other dingbat:

- Apples
- Oranges
- Bananas

A list in which each entry is composed of a set of one or more terms and an associated description:

Fruits

such as apples, oranges, and more.

Nuts

Don't eat too many; you are what you eat.

Vegetables

Potatos are spelled with care.

A list in which each entry is marked with a sequentially incremented label:

1. Step one
2. Step two

2.2. Links

Links can be used within your documents to refer to different sections and chapters or to refer to documents external to yours.

2.2.1. Internal links

Click on the Section 2 link to jump to the top of this chapter. Note the anchor at the section tag.


2.2.2. External links

Click on this (<http://www.linuxdoc.org/>) link to jump to the LDP site. Note you can use http, ftp, news and other protocols in the locator if required.

2.3. Images

Avoid diagrams if possible as this cannot be rendered in the ASCII outputs which are still needed by many around the world.

Figure 1. Graphics Test Image



Here is another variation which allows for ALT text:



Caption for the graphic goes here: This is a Green Ball.

3. Structure

A quick overview on how all parts fit together in the overall structure. An example from the Multi Disk HOWTO is used.

As this type of document is supposed to be as much for learning as a technical reference document I have rearranged the structure to this end. For the designer of a system it is more useful to have the information presented in terms of the goals of this exercise than from the point of view of the logical layer structure of the devices themselves. Nevertheless this document would not be complete without such a layer structure the computer field is so full of, so I will include it here as an introduction to how it works.

3.1. Logical structure

This is based on how each layer access each other, traditionally with the application on top and the physical layer on the bottom. It is quite useful to show the interrelationship between each of the layers used in controlling drives.

__	File structure	(/usr /tmp etc)	__
__	File system	(ext2fs, vfat etc)	__
__	Volume management	(AFS)	__
__	RAID, concatenation	(md)	__
__	Device driver	(SCSI, IDE etc)	__
__	Controller	(chip, card)	__
__	Connection	(cable, network)	__
__	Drive	(magnetic, optical etc)	__

In the above diagram both volume management and RAID and concatenation are optional layers. The 3 lower layers are in hardware. All parts are discussed at length later on in this document.

3.2. Document structure

Most users start out with a given set of hardware and some plans on what they wish to achieve and how big the system should be. This is the point of view I will adopt in this document in presenting the material, starting out with hardware, continuing with design constraints before detailing the design strategy that I have found to work well. I have used this both for my own personal computer at home, a multi purpose server at work and found it worked quite well. In addition my Japanese co-worker in this project have applied the same strategy on a server in an academic setting with similar success.

Finally at the end I have detailed some configuration tables for use in your own design. If you have any comments regarding this or notes from your own design work I would like to hear from you so this document can be upgraded.

3.3. Reading plan

As you go beyond 50 pages or so there will be a lot of text that experts and even the experienced do not need to read. Keeping in mind that we wish to care for all kinds of people in the Linux world we might have to make a reading plan. Again, an example follows from the Multi Disk HOWTO.

Although not the biggest HOWTO it is nevertheless rather big already and I have been requested to make a reading plan to make it possible to cut down on the volume.

Expert

(aka the elite). If you are familiar with Linux as well as disk drive technologies you will find most of what you need in the appendices. Additionally you are recommended to read the FAQ and the Section 12chapter.

Experienced

(aka Competent). If you are familiar with computers in general you can go straight to the chapters on Section 4 and continue from there on.

Newbie

(mostly harmless). You just have to read the whole thing. Sorry. In addition you are also recommended to read all the other disk related HOWTOs.

4. Technologies

Introduction of technology for the newbie with a few references to detailed works. Remember that not everyone has Internet access so you have to explain in sufficient details so even the newbie can get by.

5. Implementation

Now your readers should have a sufficient knowledge of what this is about and now we come to the hands on of implementing your clever scheme.

6. Maintenance

Few systems and designs are maintenance free, here you explain how to keep the system running.

7. Advanced Issues

You can get most things up and running in a quick and dirty fashion, useful for testing and getting used to how things work. For more serious use you would need to be a little more advanced. This is the place to explain it all, if applicable.

8. Further Information

A HOWTO cannot describe everything, some times the user has to venture out on the net to get more information or just updates. Here is the place to tell where and how. Again examples from the Multi Disk HOWTO, replace as needed. There is wealth of information one should go through when setting up a major system, for instance for a news or general Internet service provider. The FAQs in the following groups are useful:

8.1. News groups

Some of the most interesting news groups are:

- Storage (news:comp.arch.storage).
- PC storage (news:comp.sys.ibm.pc.hardware.storage).
- AFS (news:alt.filesystems.afs).
- SCSI (news:comp.periphs.scsi).

- Linux setup (news:comp.os.linux.setup).

Most newsgroups have their own FAQ that are designed to answer most of your questions, as the name Frequently Asked Questions indicate. Fresh versions should be posted regularly to the relevant newsgroups. If you cannot find it in your news spool you could go directly to the FAQ main archive FTP site (<ftp://rtfm.mit.edu/>). The WWW versions can be browsed at the FAQ main archive WWW site (<http://www.cis.ohio-state.edu/hypertext/faq/usenet/FAQ-List.html>).

Some FAQs have their own home site, of particular interest:

- SCSI FAQ (http://www.paranoia.com/~filipg/HTML/LINK/F_SCSI.html) and
- comp.arch.storage FAQ (http://alumni.caltech.edu/~rdv/comp_arch_storage/FAQ-1.html).

8.2. Mailing Lists

These are low-noise channels mainly for developers. Think twice before asking questions there as noise delays the development. Some relevant lists are `<linux-raid>`, `<linux-scsi>` and `<linux-ext2fs>`. Many of the most useful mailing lists run on the `vger.rutgers.edu` server but this is notoriously overloaded, so try to find a mirror. There are some lists mirrored at The Redhat Home Page (<http://www.redhat.com>). Many lists are also accessible at linuxhq (<http://www.linuxhq.com/lxlists>), and the rest of the web site contains useful information as well.

If you want to find out more about the lists available you can send a message with the line **lists** to the list server at `<majordomo@vger.rutgers.edu>`. If you need help on how to use the mail server just send the line **help** to the same address. Due to the popularity of this server it is likely it takes a bit of time before you get a reply or even get messages after you send a **subscribe** command.

There is also a number of other majordomo list servers that can be of interest such as the EATA driver list (`<linux-eata@mail.uni-mainz.de>`) and the Intelligent IO list `<linux-i2o@dpt.com>`.

Mailing lists are in a state of flux but you can find links to a number of interesting lists from the Linux Documentation Homepage (<http://www.linuxdoc.org/>).

8.3. HOWTO

These are intended as the primary starting points to get the background information as well as show you how to solve a specific problem. Some relevant HOWTOs are `Bootdisk`, `Installation`, `SCSI` and `UMSDOS`. The main site for these is the LDP archive (<http://www.linuxdoc.org/>) at Metalab (formerly known as Sunsite).

There is a new HOWTO out that deals with setting up a DPT RAID system, check out the DPT RAID HOWTO homepage (http://www.ram.org/computing/linux/dpt_raid.html).

8.4. Local Resources

In most distributions of Linux there is a document directory installed, have a look in the `/usr/doc` directory. where most packages store their main documentation and README files etc. Also you will here find the HOWTO archive (`/usr/doc/HOWTO`) of ready formatted HOWTOs and also the mini-HOWTO archive (`/usr/doc/HOWTO/mini`) of plain text documents.

Many of the configuration files mentioned earlier can be found in the `/etc` directory. In particular you will want to work with the `/etc/fstab` file that sets up the mounting of partitions and possibly also `/etc/raidtab` file that is used for the `md` system to set up RAID.

The kernel source in `/usr/src/linux` is, of course, the ultimate documentation. In other words, “use the source, Luke”. It should also be pointed out that the kernel comes not only with source code which is even commented (well, partially at least) but also an informative `/usr/src/linux/Documentation`. If you are about to ask any questions about the kernel you should read this first, it will save you and many others a lot of time and possibly embarrassment.

Also have a look in your system log file (`/var/log/messages`) to see what is going on and in particular how the booting went if too much scrolled off your screen. Using **tail -f /var/log/messages** in a separate window or screen will give you a continuous update of what is going on in your system.

You can also take advantage of the `/proc` file system that is a window into the inner workings of your system. Use **cat** rather than **more** to view the files as they are reported as being zero length. Reports are that **less** works well here.

8.5. Web Sites

There are a huge number of informative web sites available. By their very nature they change quickly so do not be surprised if these links become quickly outdated.

A good starting point is of course the Linux Documentation Project (<http://www.linuxdoc.org/>) home page, an information central for documentation, project pages and much more.

Please let me know if you have any other leads that can be of interest.

9. Getting Help

In the end you might find yourself unable to solve your problems and need help from someone else. The most efficient way is either to ask someone local or in your nearest Linux user group, search the web for the nearest one.

Another possibility is to ask on Usenet News in one of the many, many newsgroups available. The problem is that these have such a high volume and noise (called low signal-to-noise ratio) that your question can easily fall through unanswered.

No matter where you ask it is important to ask well or you will not be taken seriously. Saying just *my disk does not work* is not going to help you and instead the noise level is increased even further and if you are lucky someone will ask you to clarify.

Instead describe your problems in some detail that will enable people to help you. The problem could lie somewhere you did not expect. Therefore you are advised to list the following information about your system:

Hardware

- Processor
- DMA
- IRQ
- Chip set (LX, BX etc)
- Bus (ISA, VESA, PCI etc)
- Expansion cards used (Disk controllers, video, IO etc.)

Software

- BIOS (On motherboard and possibly SCSI host adapters)
- LILO, if used
- Linux kernel version as well as possible modifications and patches
- Kernel parameters, if any
- Software that shows the error (with version number or date)

Peripherals

- Type of disk drives with manufacturer name, version and type

- Other relevant peripherals

Remember that booting text is logged to `/var/log/messages` which can answer most of the questions above. Obviously if the drives fail you might not be able to get the log saved to disk but you can at least scroll back up the screen using the **SHIFT** and **PAGE UP** keys. It may also be useful to include part of this in your request for help but do not go overboard, keep it *brief* as a complete log file dumped to Usenet News is more than a little annoying.

10. Concluding Remarks

Just summing up... Also a place for general recommendations.

11. Questions and Answers

Check the newsgroups and try to determine some frequent problems and cover them here. Again an example from the Multi Disk HOWTO.

This is just a collection of what I believe are the most common questions people might have. Give me more feedback and I will turn this section into a proper FAQ.

- Q:How many physical disk drives (spindles) does a Linux system need?

A: Linux can run just fine on one drive (spindle). Having enough RAM (around 32 MB, and up to 64 MB) to support swapping is a better price/performance choice than getting a second disk. (E)IDE disk is usually cheaper (but a little slower) than SCSI.

- Q: Are there any disadvantages in this scheme?

A: There is only a minor snag: if even a single partition overflows the system might stop working properly. The severity depends of course on what partition is affected. Still this is not hard to monitor, the command **df** gives you a good overview of the situation. Also check the swap partition(s) using **free** to make sure you are not about to run out of virtual memory.

- Q: OK, so should I split the system into as many partitions as possible for a single drive?

A: No, there are several disadvantages to that. First of all maintenance becomes needlessly complex and you gain very little in this. In fact if your partitions are too big you will seek across larger areas than needed. This is a balance and dependent on the number of physical drives you have.

* *Greg Leblanc: Depending on how big this FAQ gets, perhaps it would be worthwhile to have, say, the 5 most FAQ, and put the rest into an external FAQ. Dunno. Comments?*
(rest deleted.)

12. Bits and Pieces

This is basically a section where I stuff all the bits I have not yet decided where should go, yet that I feel is worth knowing about. It is a kind of transient area.

13. Examples

Example designs and sample configuration files and other relevant details is always handy